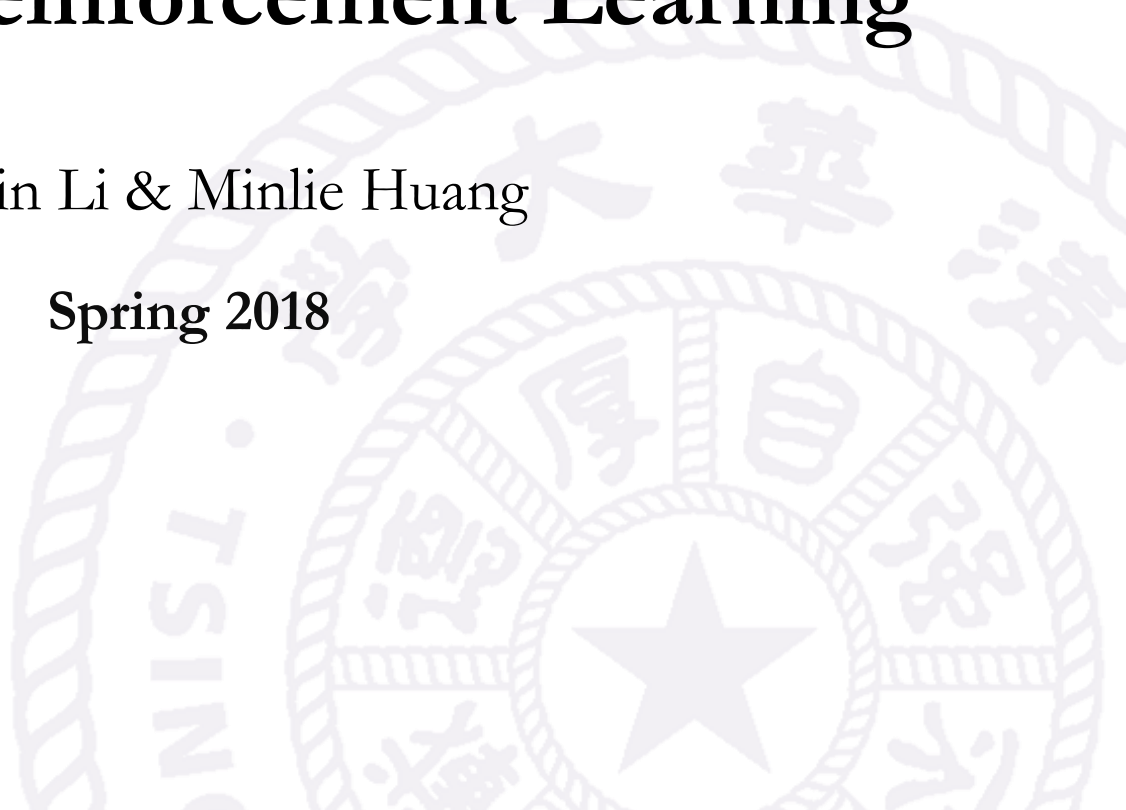


Introduction to Artificial Intelligence

Project 4 – Reinforcement Learning

Jianmin Li & Minlie Huang

Spring 2018




Reinforcement Learning and Pacman

- Berkeley RL Pacman Project
 - <http://ai.berkeley.edu/reinforcement.html>
 - <https://s3-us-west-2.amazonaws.com/cs188websitecontent/projects/release/reinforcement/v1/001/reinforcement.zip>



Get Started

- `python gridworld.py -m`
- `python gridworld.py -g MazeGrid`



Basic Tasks

- Value Iteration
 - Implement **ValueIterationAgent** in **valueIterationAgents.py**
 - `python gridworld.py -a value -i 100 -k 10`
 - **7 points**
- Parameter Analysis in Value Iteration
 - Change parameters so that the optimal policy causes the agent to attempt to cross the bridge
 - Put your answer in **question2()** of **analysis.py**
 - `python gridworld.py -a value -i 100 -g BridgeGrid --discount 0.9 --noise 0.2`
 - **1 point**

Basic Tasks

- Q-Learning
 - implement the `update`, `computeValueFromQValues`, `getQValue`, `computeActionFromQValues` methods in `QLearningAgent` in `qlearningAgents.py`
 - `python gridworld.py -a q -k 5 -m`
 - 6 points
- Q-Learning and Pacman
 - Play pacman with Q-Learning
 - `python pacman.py -p PacmanQAgent -x 2000 -n 2010 -l smallGrid`
 - `QLearningAgent.getAction` is provided. Some code in `QLearningAgent` may be modified
 - 1 point

Bonus

- **Approximate Q-Learning**

- Implement **ApproximateQAgent** in **qlearningAgents.py**
- `python pacman.py -p ApproximateQAgent -x 2000 -n 2010 -l smallGrid`
- **1 point**
- NOTE: Deep QL is inspired from this



Submission

- A 2-3 pages report (either Chinese or English)
 - Compare with different parameters in Value Iteration and Q-Learning, etc.
- Zip the files as the following structure
 - student_id.zip (e.g. 20090112xx.zip)
 - student_id.pdf
 - valueIterationAgents.py
 - qlearningAgents.py
 - analysis.py



Grading

- Due
 - 2018/7/4 23:59:59
 - For Senior/Graduate Students: 2018/6/25 23:59:59
- Correctness of algorithms (80%)
- Report (20%)

